

ENGINEERING KIOSK ALPS | 12.06.2025 | INNSBRUCK

# Multi-Tenant architectures

**Maximilian Schellhorn**

Sr. Solutions Architect  
Amazon Web Services (AWS)



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.



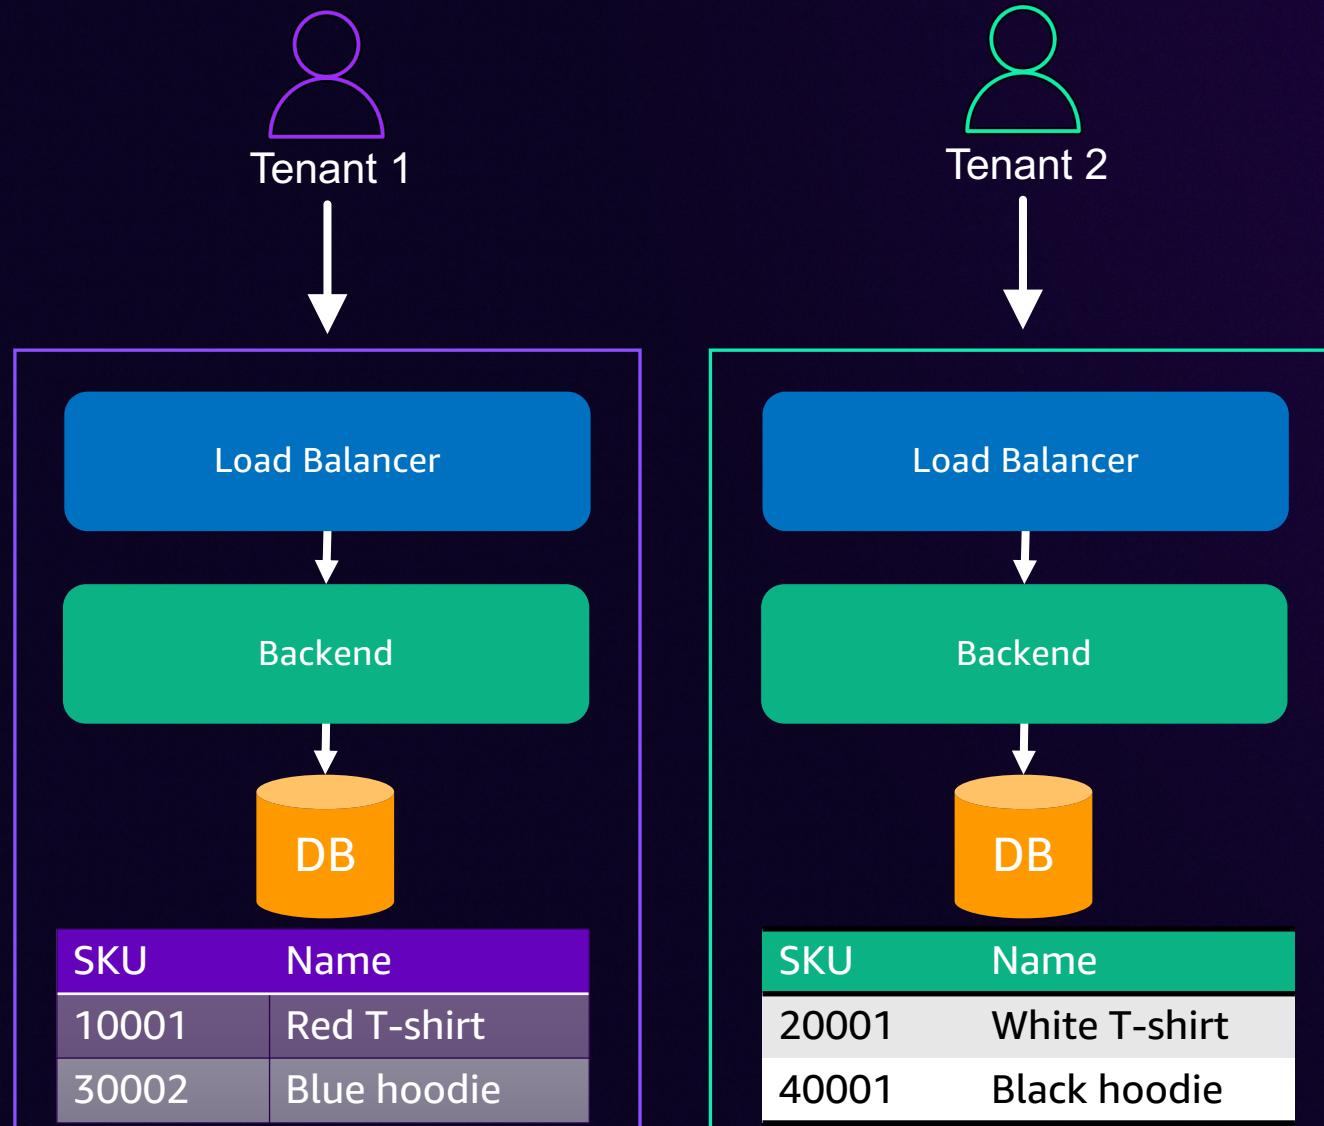


# Tools NOT Rules



# Challenge 1: One size fits all mindset

# Single-Tenant application



## Benefits

- Simple programming model
- Tenant identity & isolation by design
- Smaller blast radius

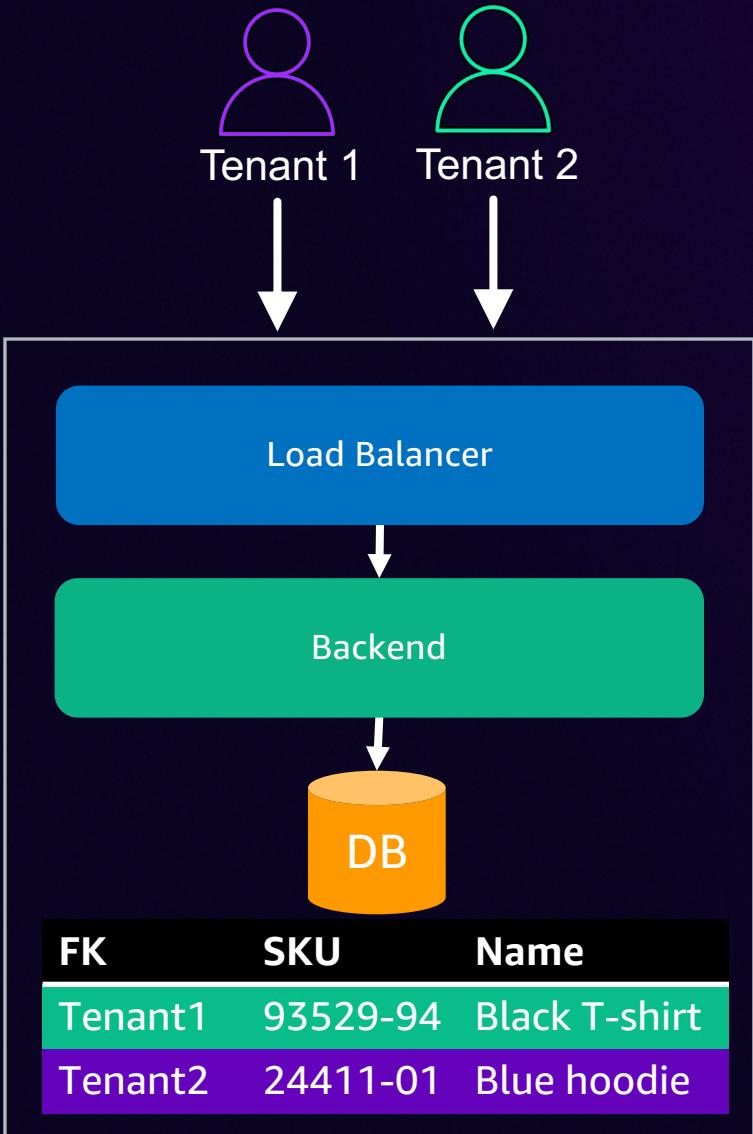
## Challenges

- Maintenance & Onboarding
- Scales linear (time, effort, money)
- Enables version drift

## Use cases

- Static set of (high paying) customers
- Capped amount of potential customers
- Conservative/Air gapped customers only

# Multi-Tenant application



## Benefits

Efficient & scalable  
Less maintenance & faster updates

## Challenges

More complex programming model  
Tenant identity & isolation at runtime  
Blast radius & Noisy neighbors

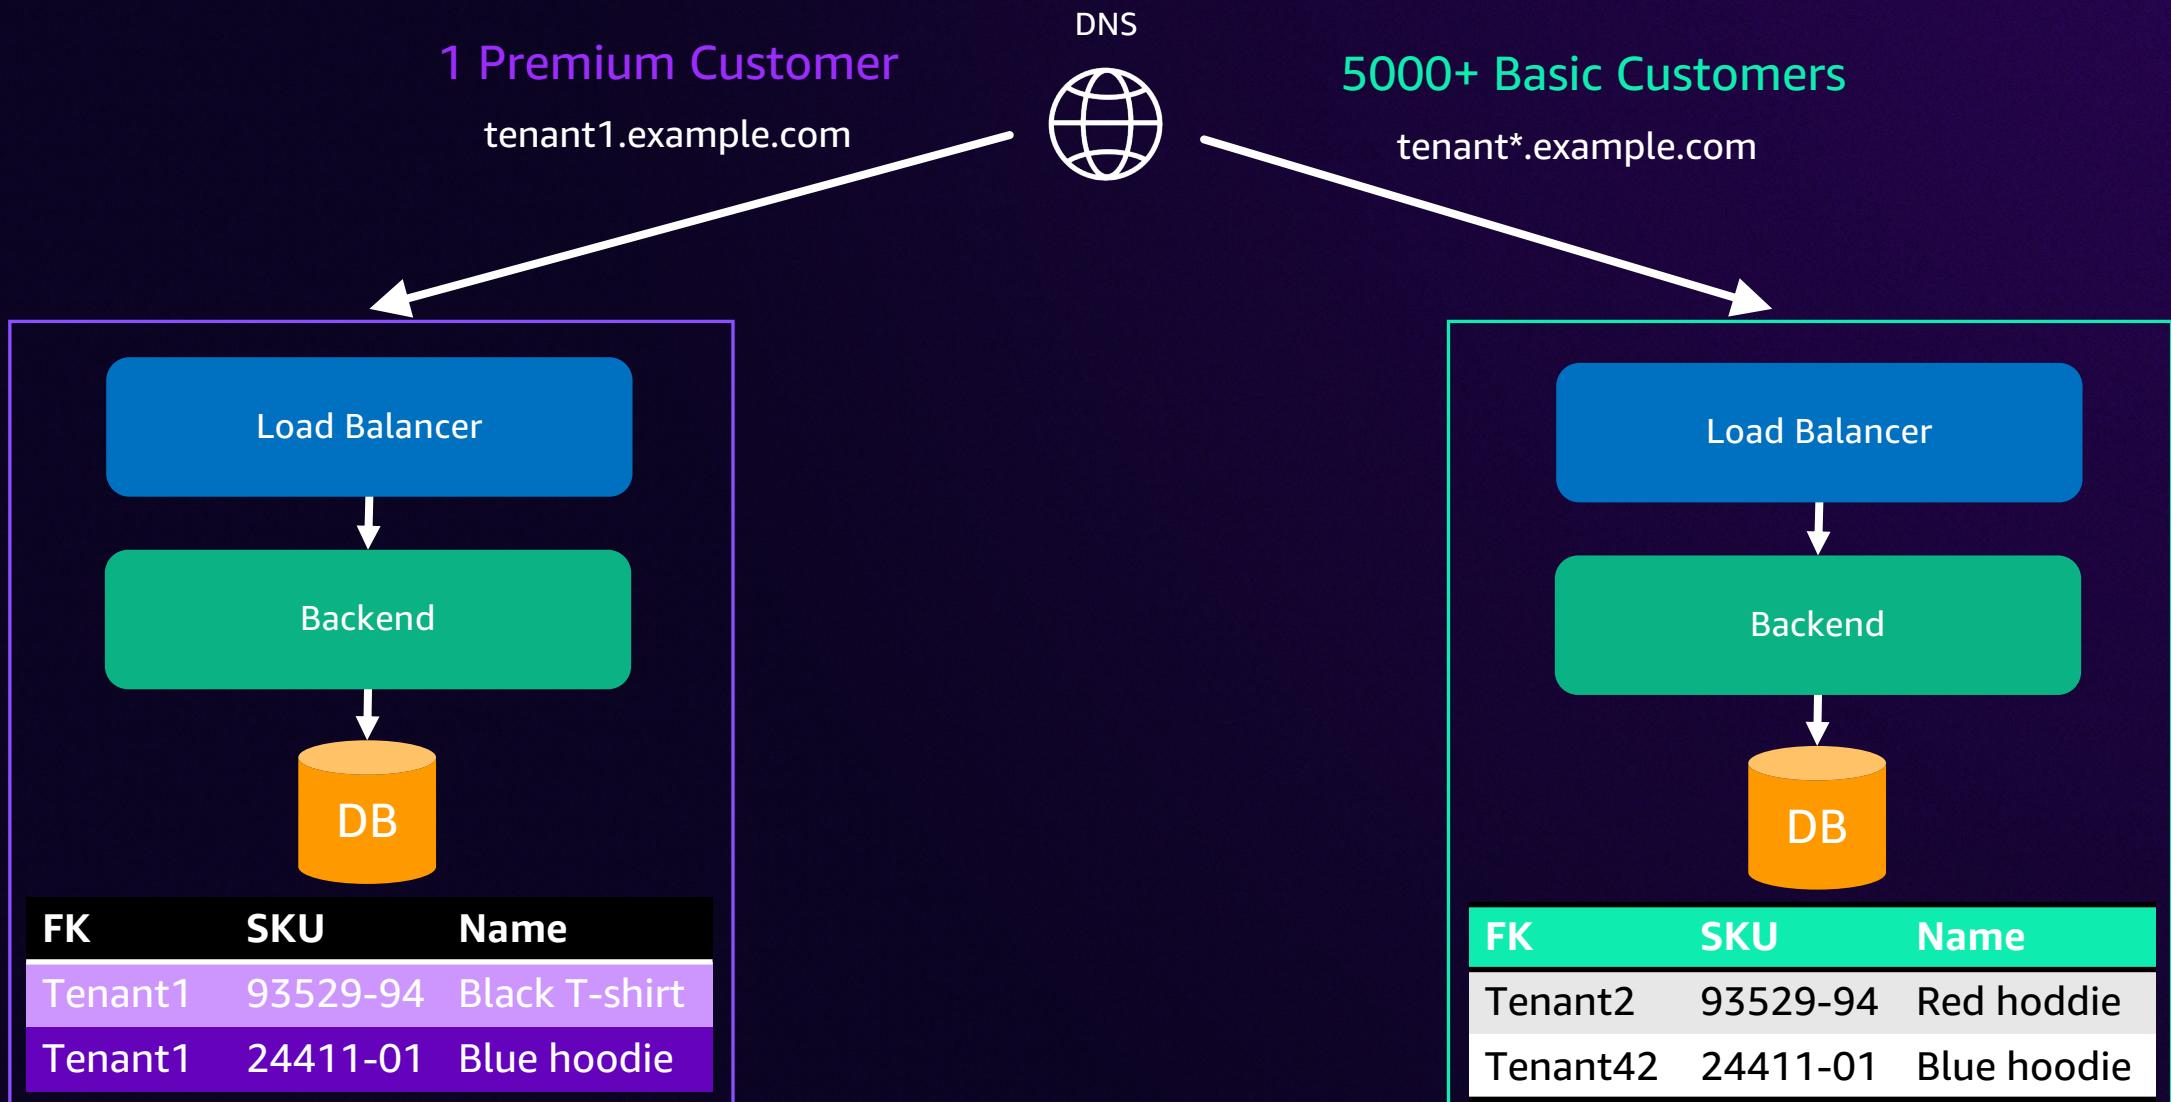
## Use cases

Rapid scale & customer growth  
Fast time to market & onboarding  
Software-as-a-Service (SaaS)

# **Multi-Tenant systems are better Single-Tenant systems**

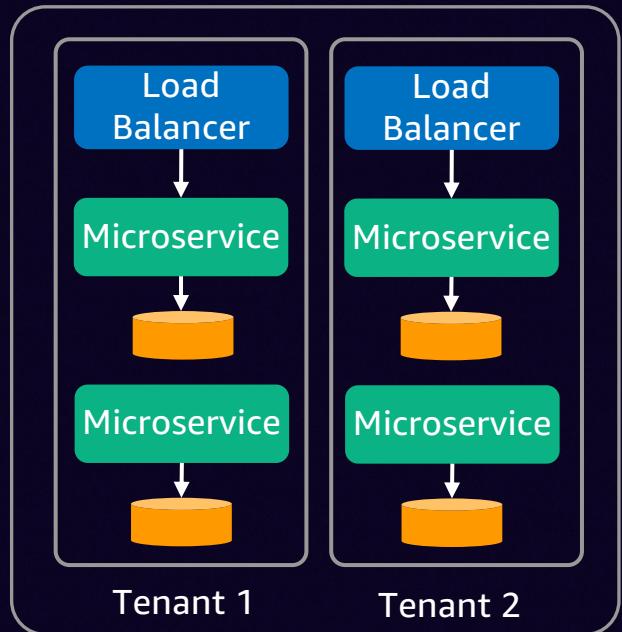
A Multi-Tenant system deployed for 1 customer = Single-Tenant semantics

# Multi-Tenant applications provide **flexibility**

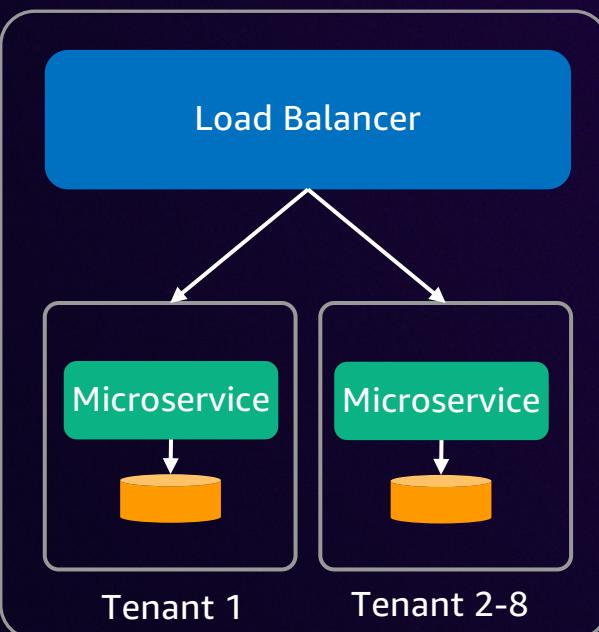


# Multi-Tenant deployment models

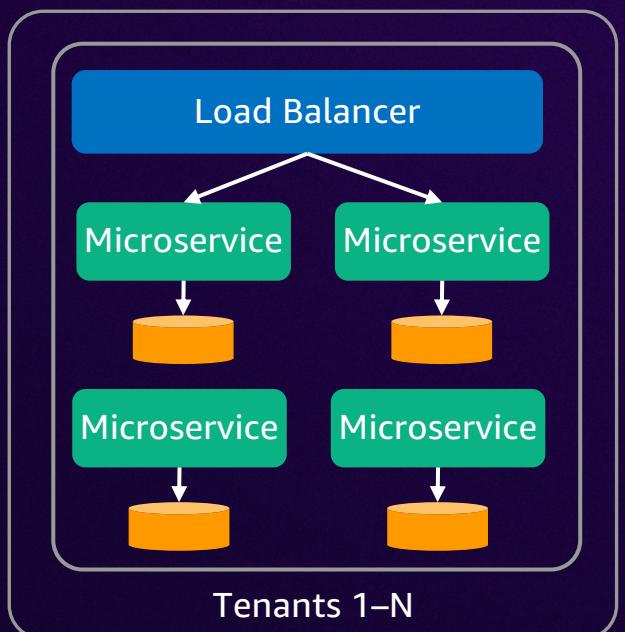
Silo



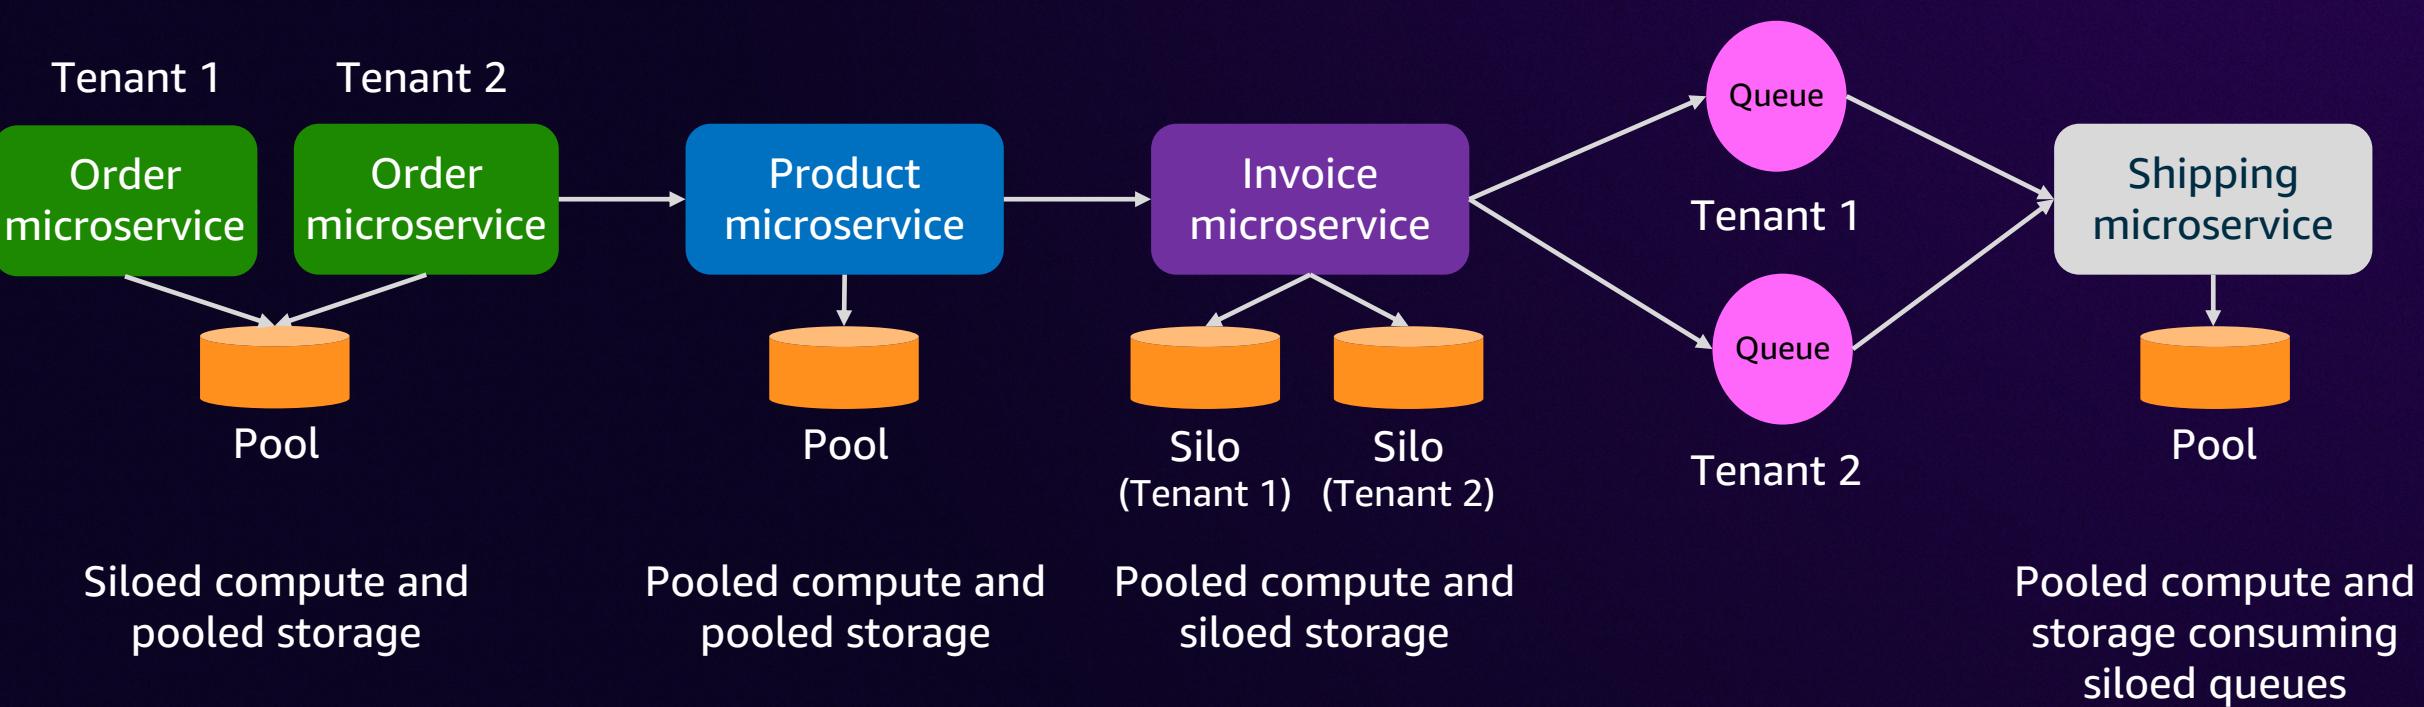
Bridge



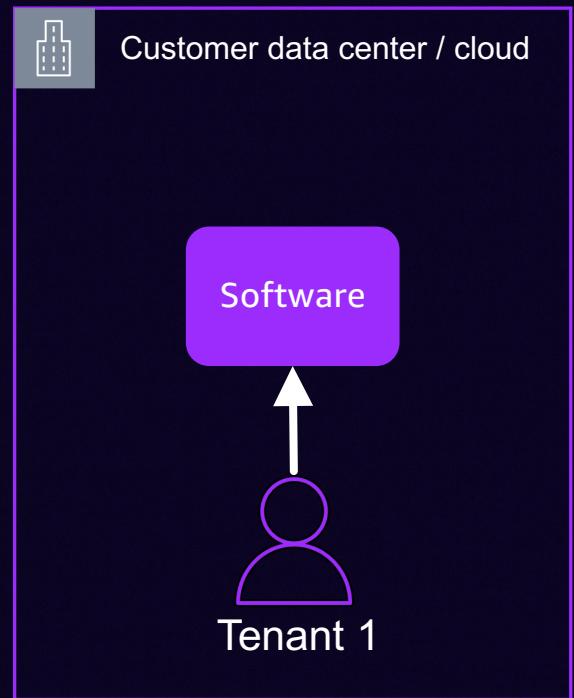
Pool



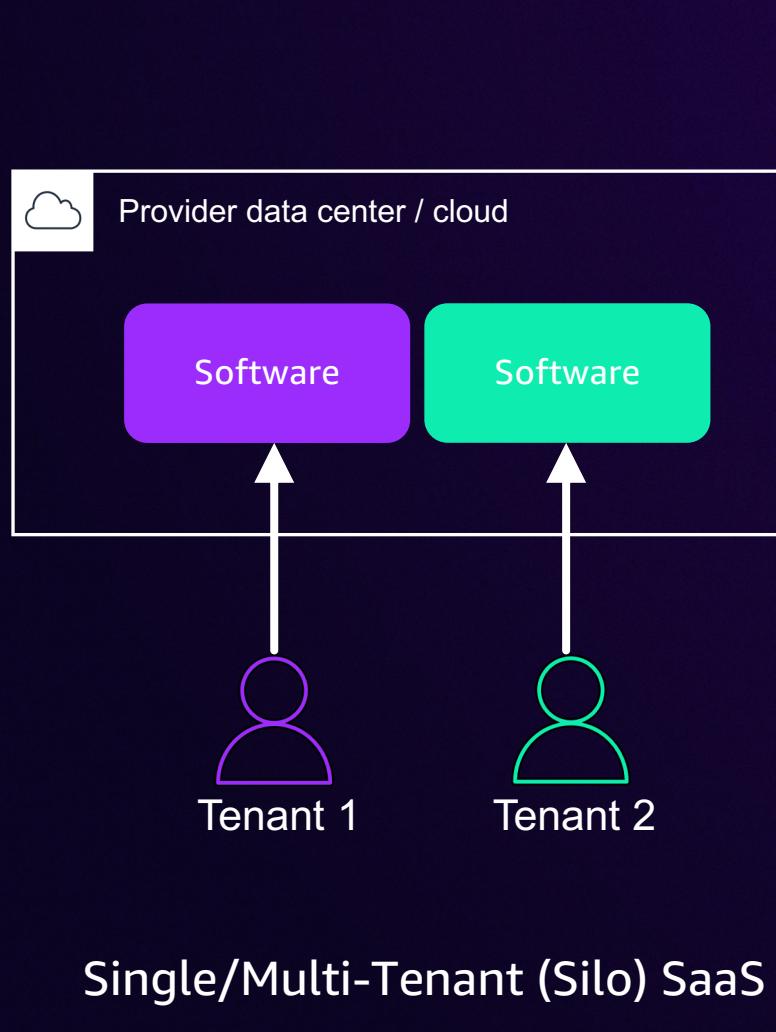
# Multi-Tenant deployments in reality



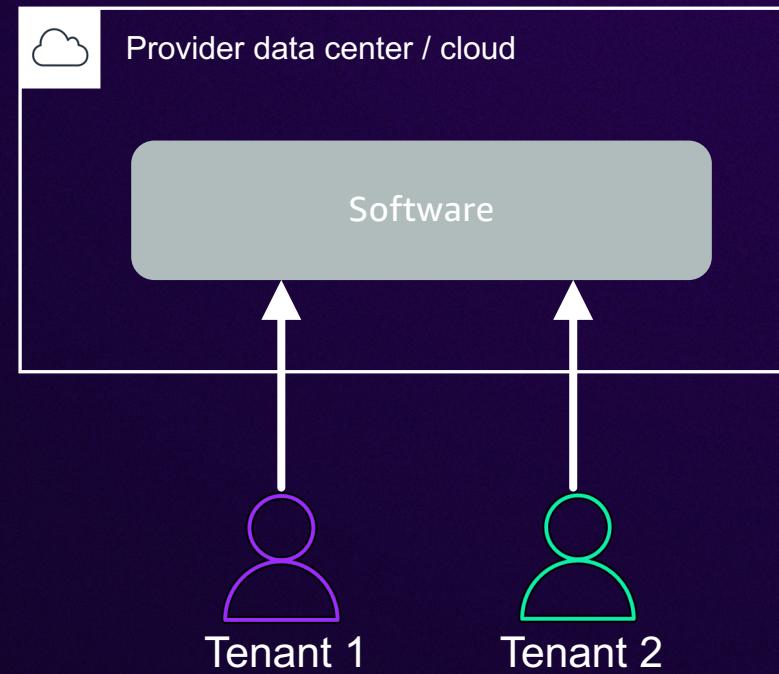
# Self-Managed vs. SaaS



Single-Tenant Self-Managed



Single/Multi-Tenant (Silo) SaaS



Multi-Tenant (Pooled) SaaS

“

Before we were doing SaaS, we did not prioritize the infrastructure footprint of our software because it was the customer's responsibility to provide and pay for the hardware. **With SaaS, every dollar of unoptimized infrastructure is lost business.**

**CTO, SaaS company**

“

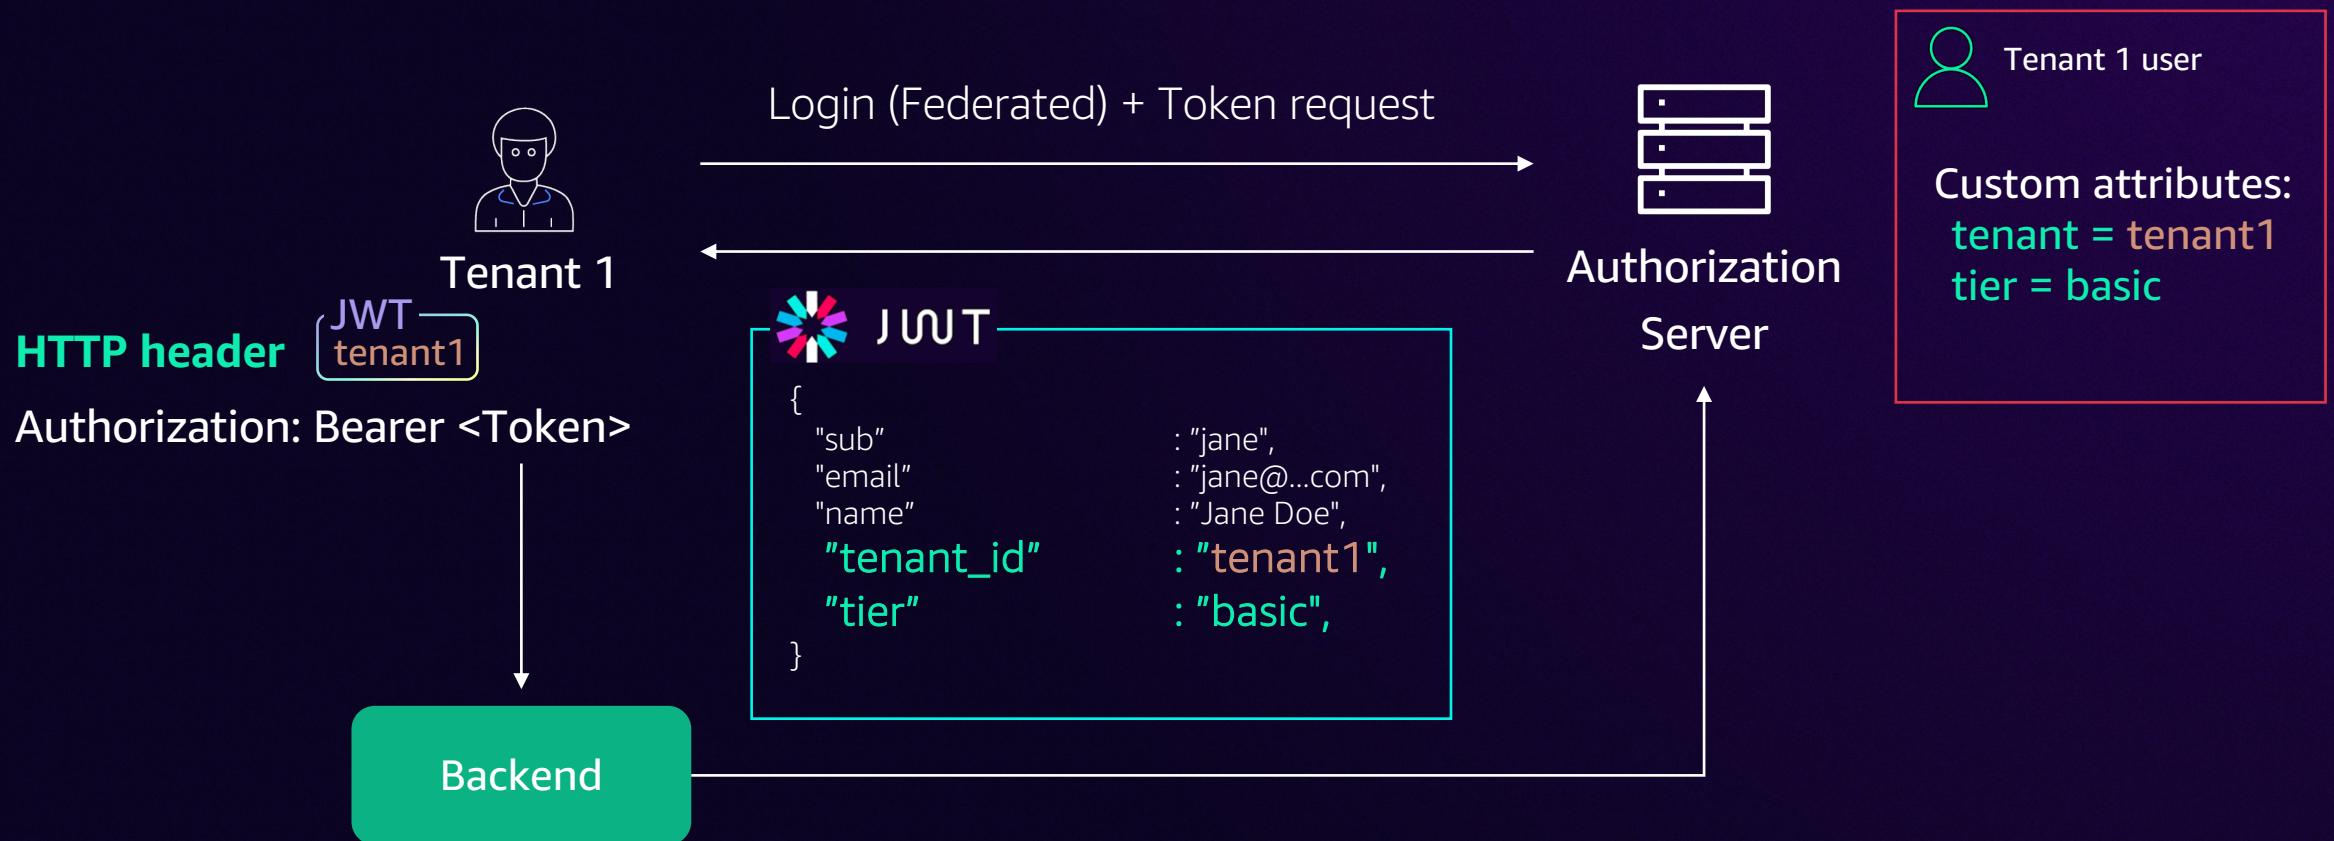
Engineering is the art of **not-constructing** rather than constructing. It is the art of doing **that well with 1\$**, which any bungler can do with 2\$ after fashion.

**Arthur M. Wellington**

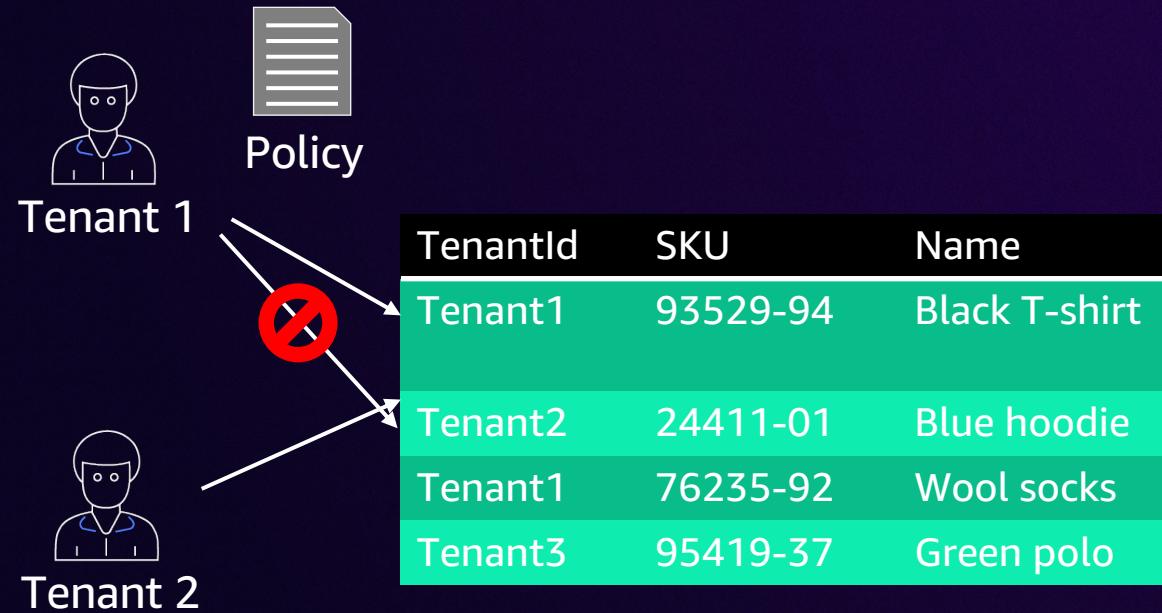


# Challenge 2: Tenant identity & isolation

# Runtime identity with tokens (JWT)

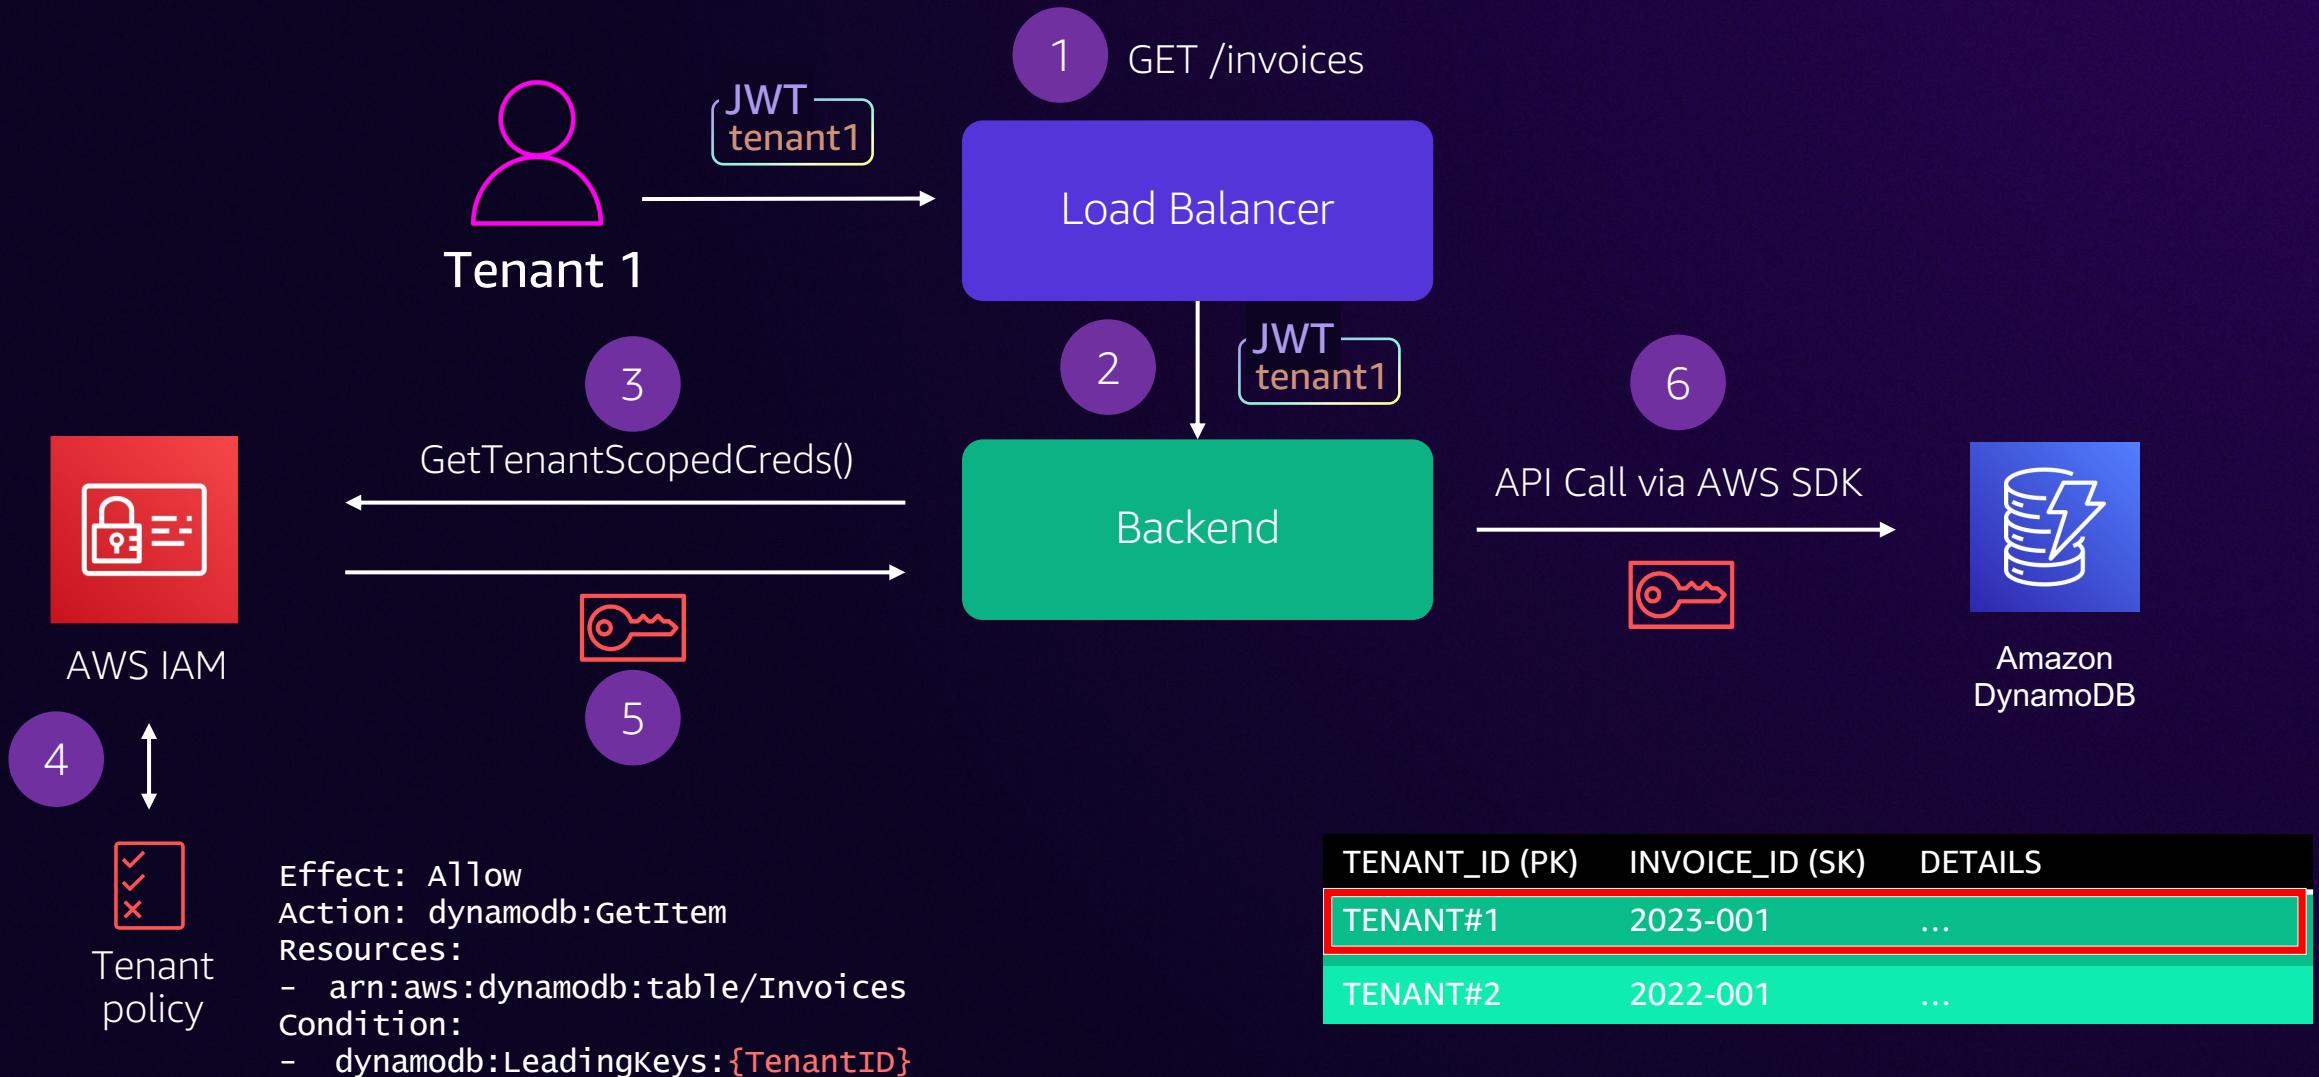


# Tenant isolation at runtime



Single table separated by identifier

# API based runtime isolation



# PostgreSQL Row Level Security (RLS)

## Initialize RLS

```
-- Turn on RLS
ALTER TABLE tenant ENABLE ROW LEVEL SECURITY;

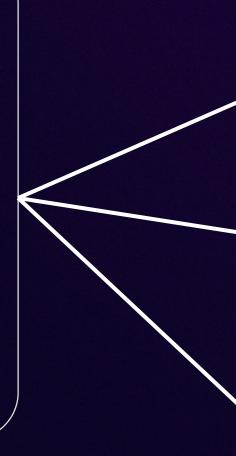
-- Scope read/write by tenant
CREATE POLICY tenant_isolation_policy ON tenant
USING (tenant_id = (current_setting('app.current_tenant')::text));
```

## Query with RLS

```
-- SET tenant context
rls_multi_tenant=> SET app.current_tenant = 'tenant1';

-- No tenant context required
rls_multi_tenant=> SELECT * FROM tenant;

-- Attempt to force other tenant id would not work
rls_multi_tenant=> SELECT * FROM tenant WHERE tenant_id = 'tenant2'
```



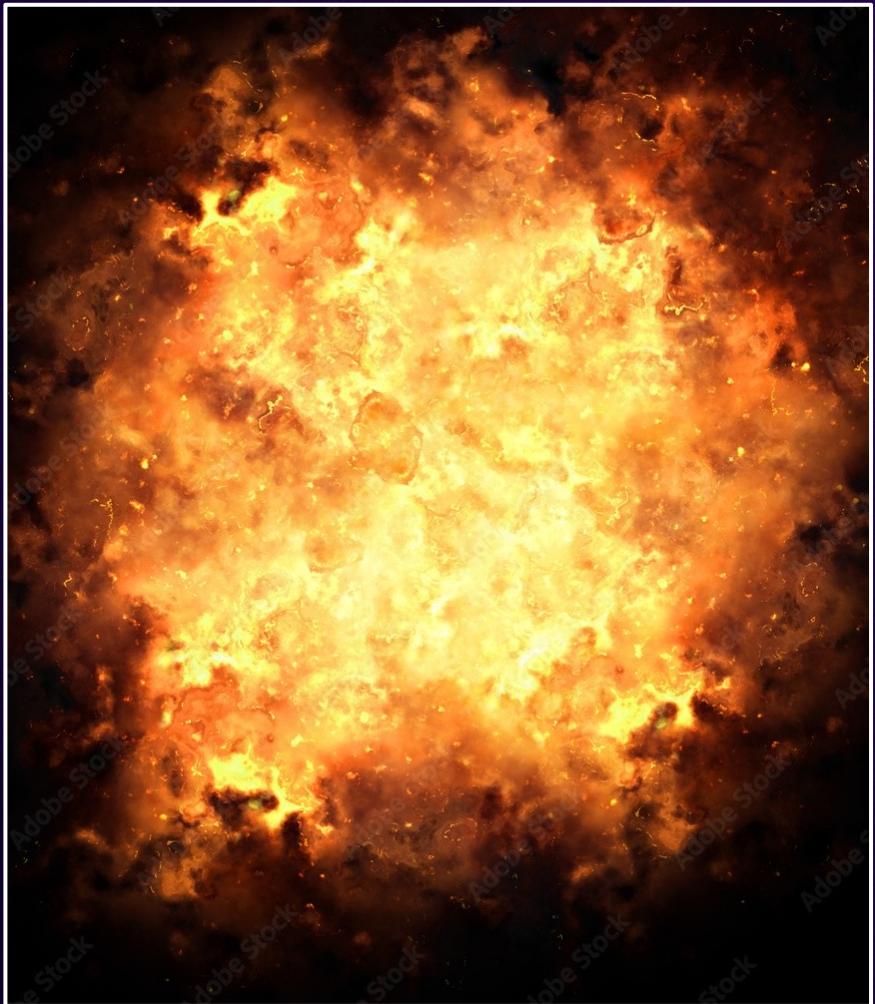
FK	SKU	Name
Tenant1	93529-94	Black T-shirt
Tenant2	24411-01	Blue hoodie
Tenant1	76235-92	Wool socks
Tenant3	95419-37	Green polo
Tenant2	88314-99	White hat
Tenant1	24598-72	Tennis shoes

# Challenge 3: Blast radius & noisy neighbors

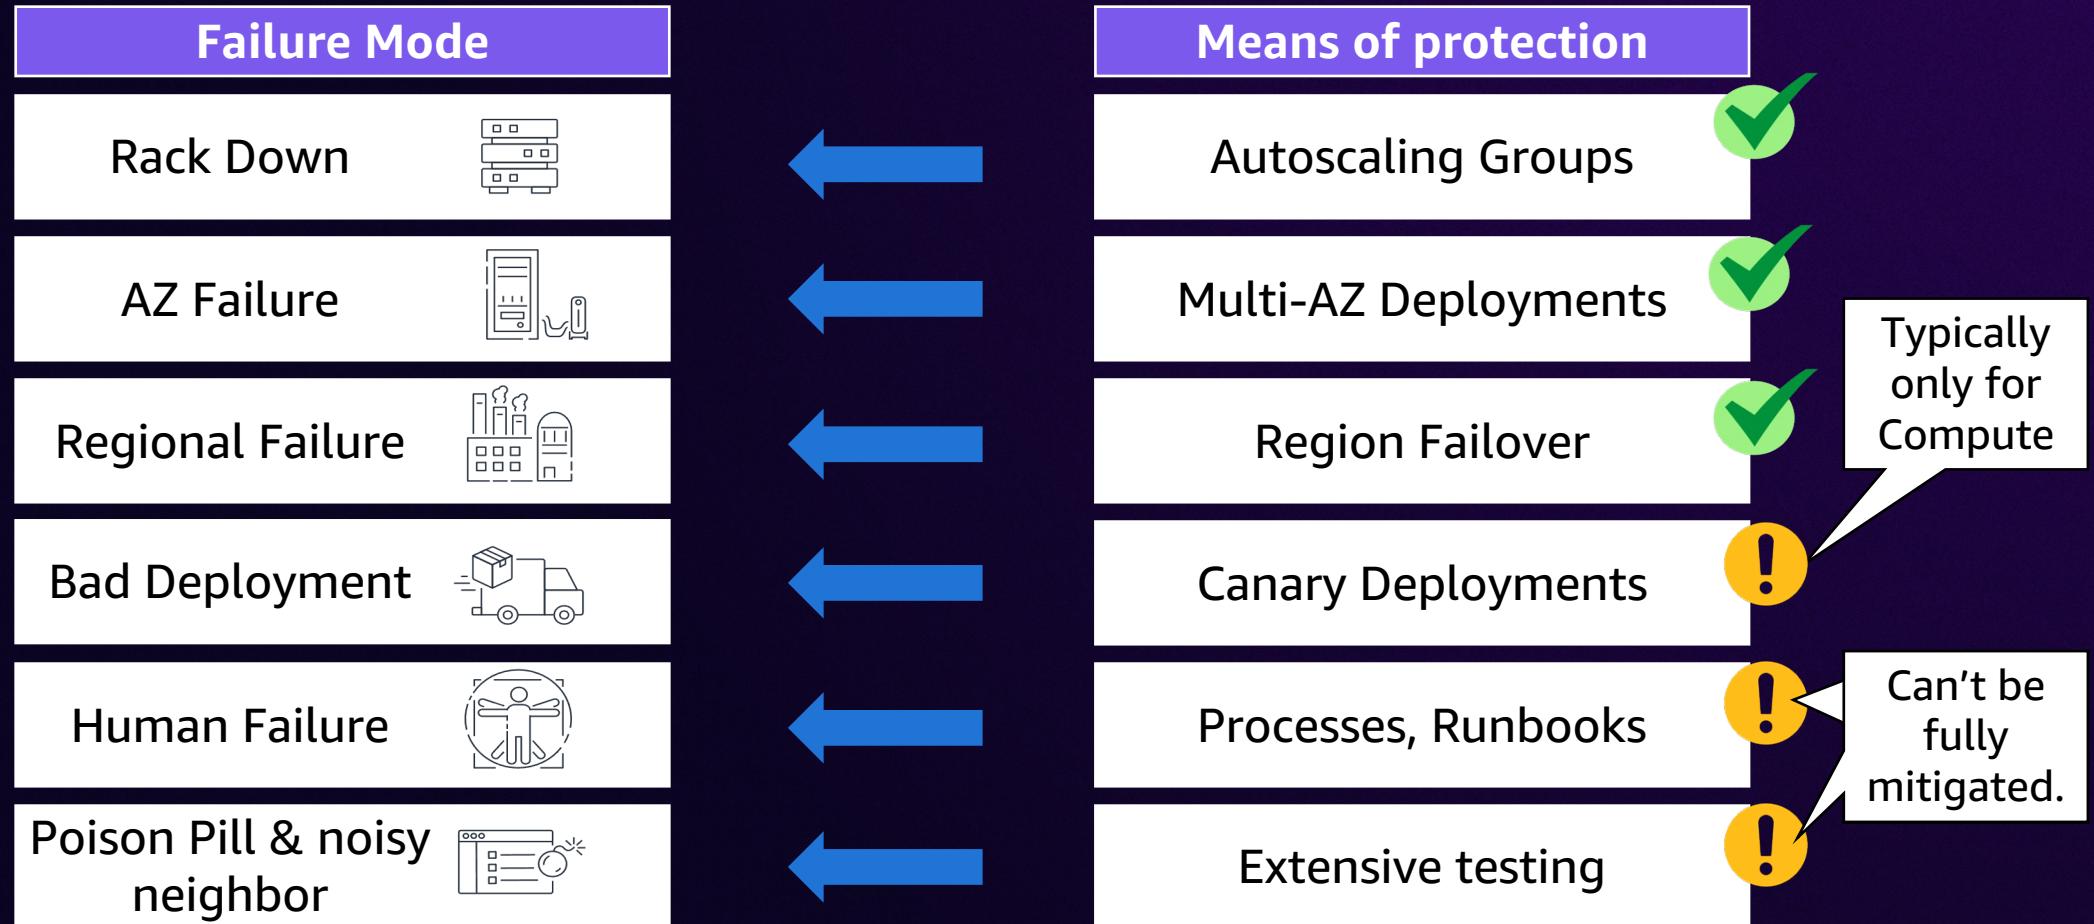
# Ignoring your **blast radius**

*"The maximum impact that might be affected in the event of a system failure"*

0 - 100%



# How do we usually protect against failures?



# But what about ...

*A tenant triggering a very specific bug?*

*A deployment that makes the production network unusable?*

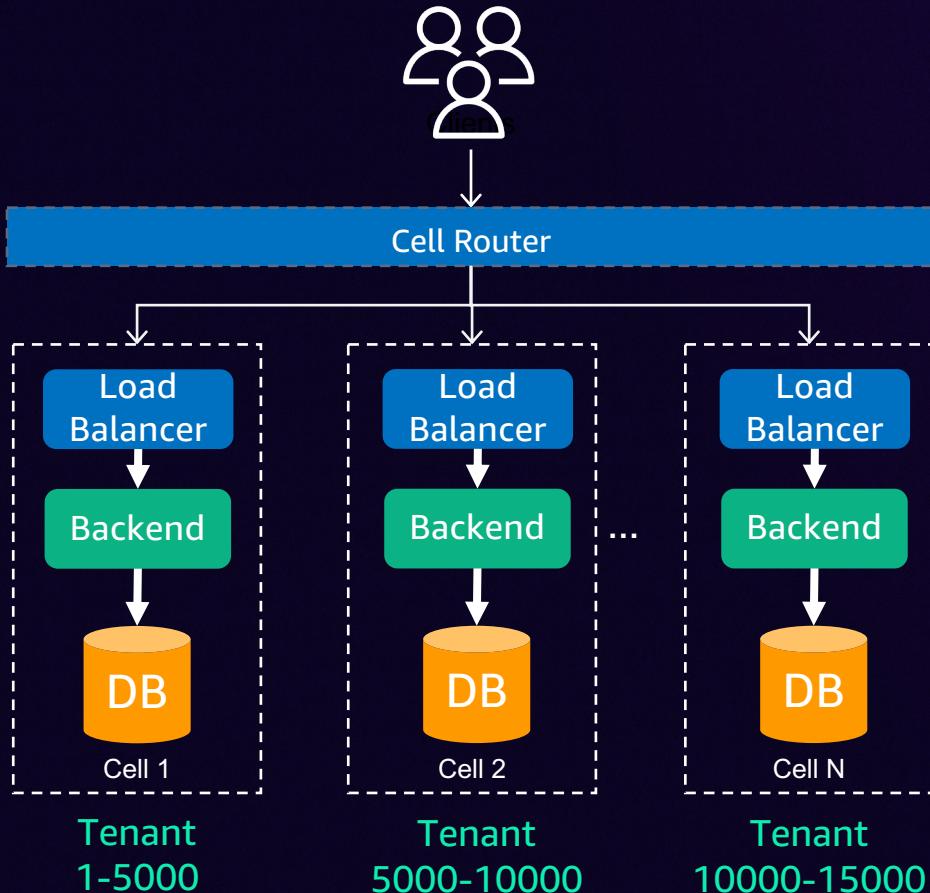
*A poisoned request that wipes the database?*

*A spike from one customer saturating all capacity?*

*Someone accidentally deleting the main loadbalancer?*

Testing makes these failure unlikely to happen.  
At scale, even unlikely scenarios happen regularly.

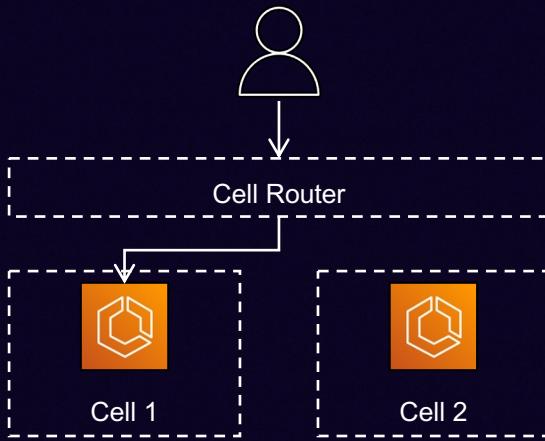
# Solution: Cell based architecture



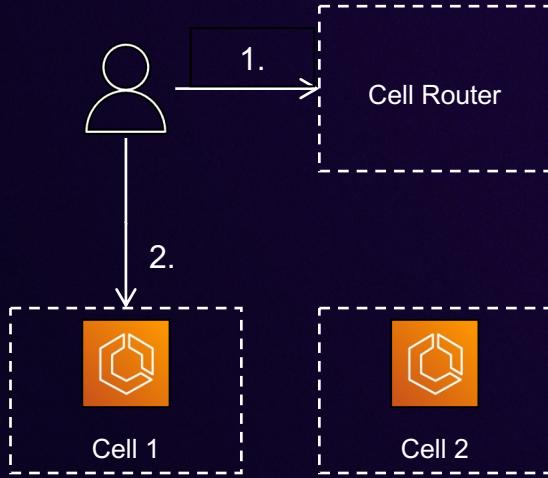
- Multiple copies of the **entire application** deployed in each cell.
- Data is **partitioned** – there is no replication between cells.
- **Complete isolation** between cells limit and contain failure.
- **Linear scale out** for additional clients
- Testability & Rollout benefits

# Cell-Routing options

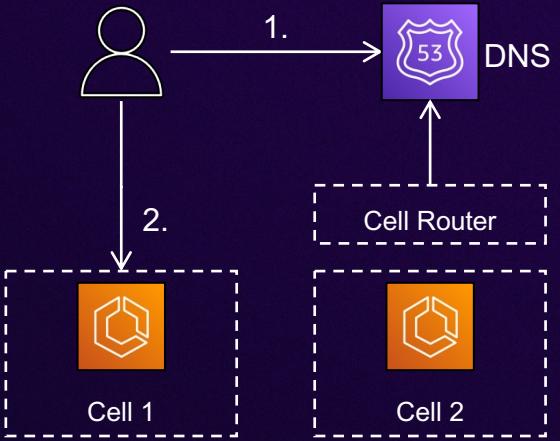
## A) Cell Router as Load Balancer



## B) Router forwards the request



## C) Routing through DNS



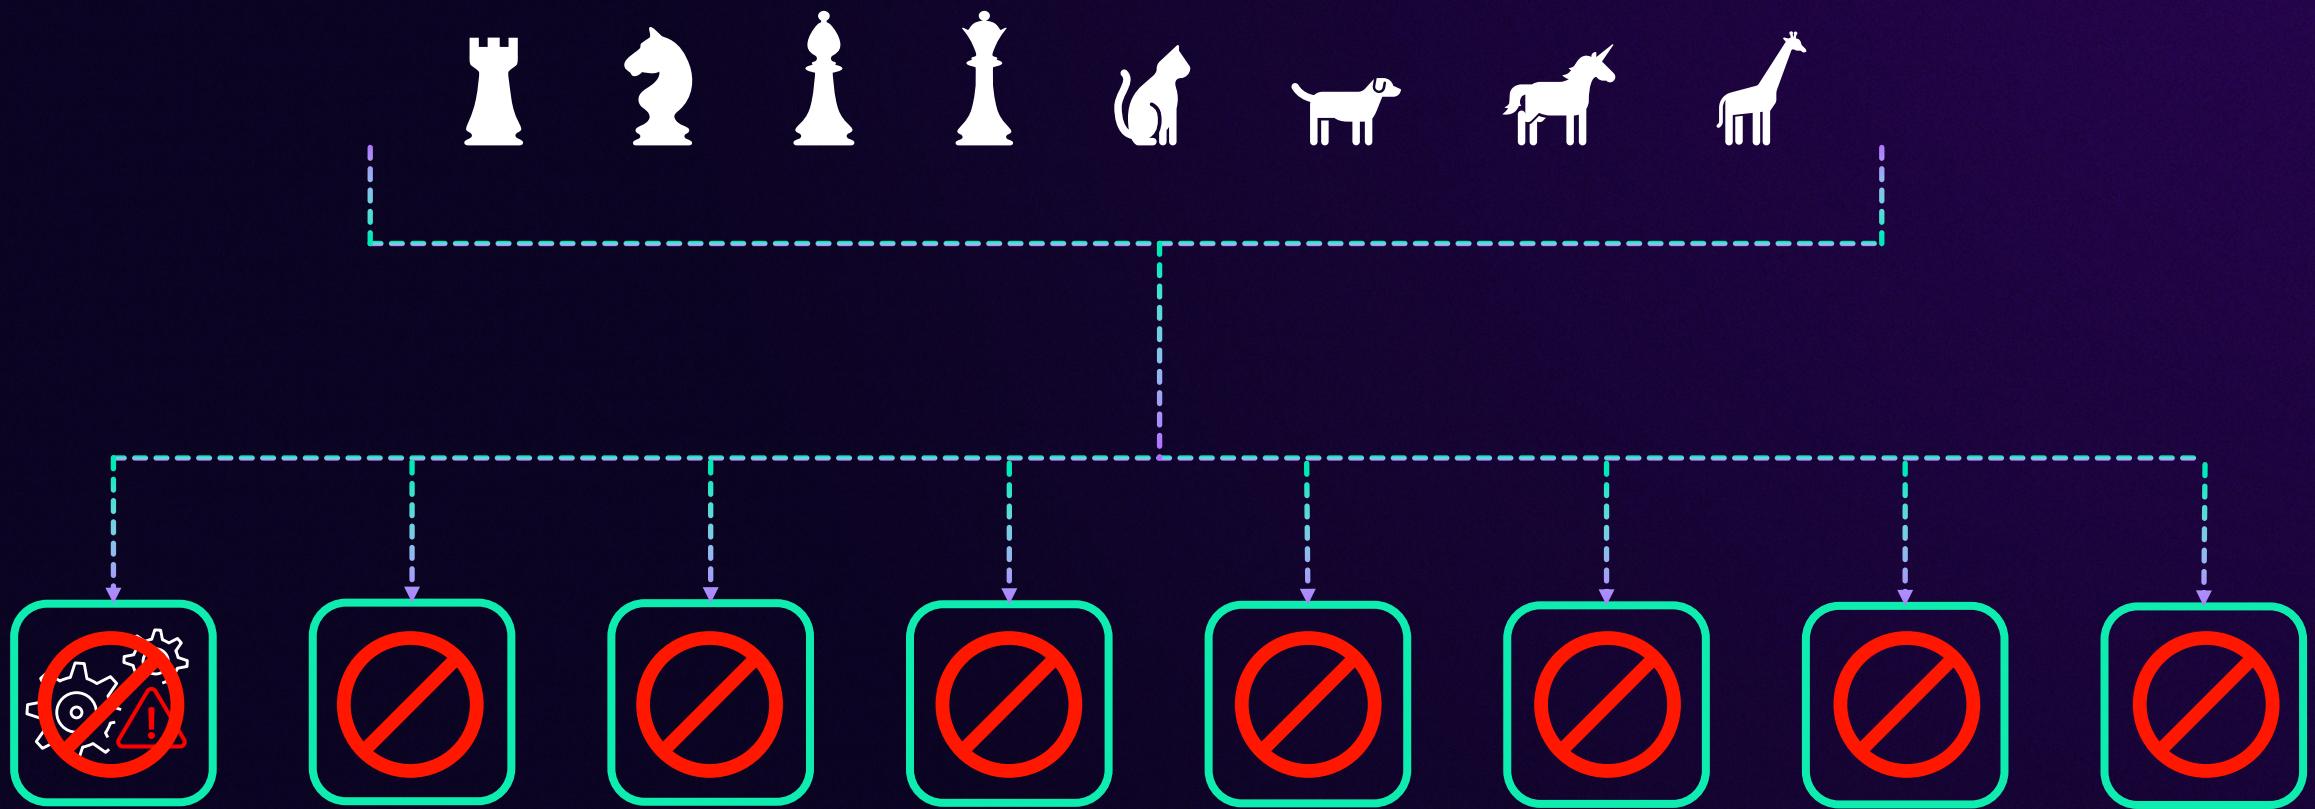
- ✓ Transparent to clients
- Cell Router is critical to ongoing transactions

- ✓ Simple to implement
- ✓ Clients with cached routing information unaffected by failure
- Custom logic needed on client
- Potentially higher latency

- ✓ Most DNS providers offer high availability.
- Clients need to map users to DNS names

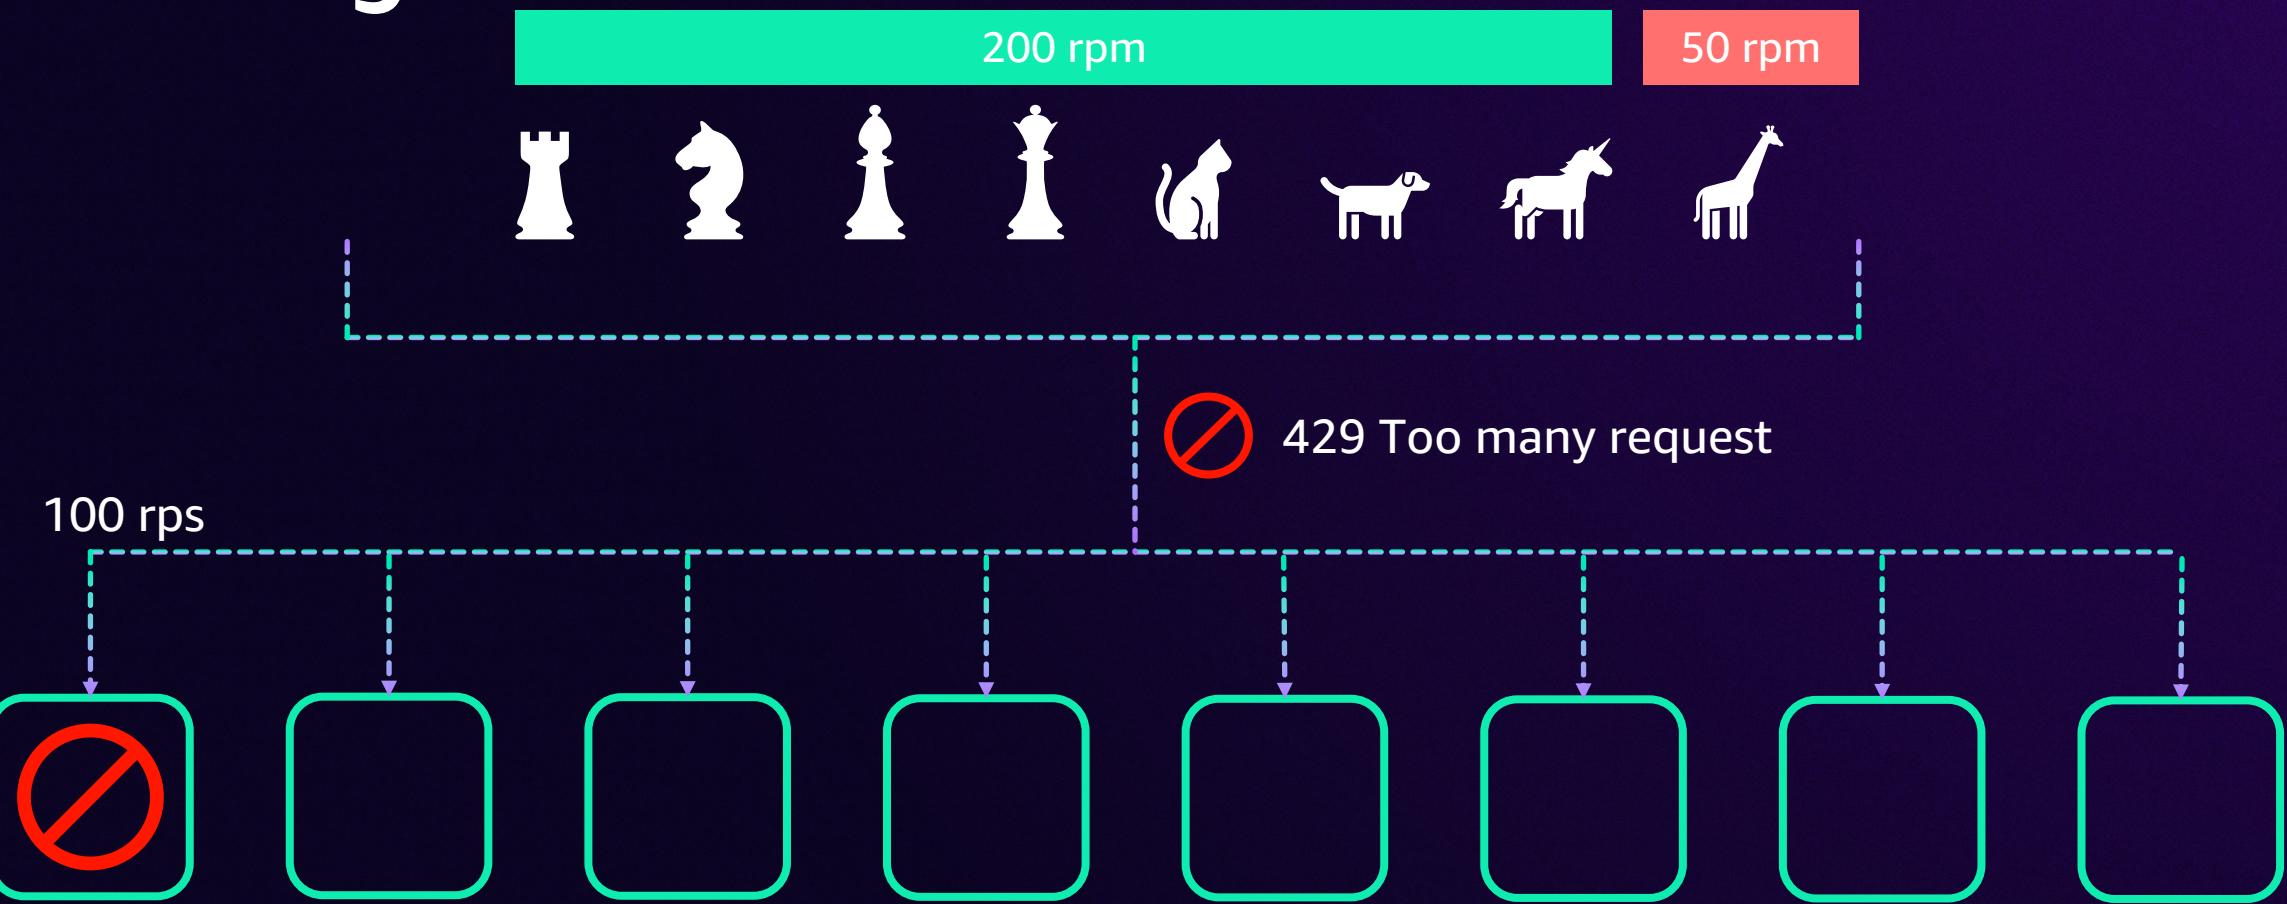
# Noisy Neighbor & Poison pills

# Traditional architecture

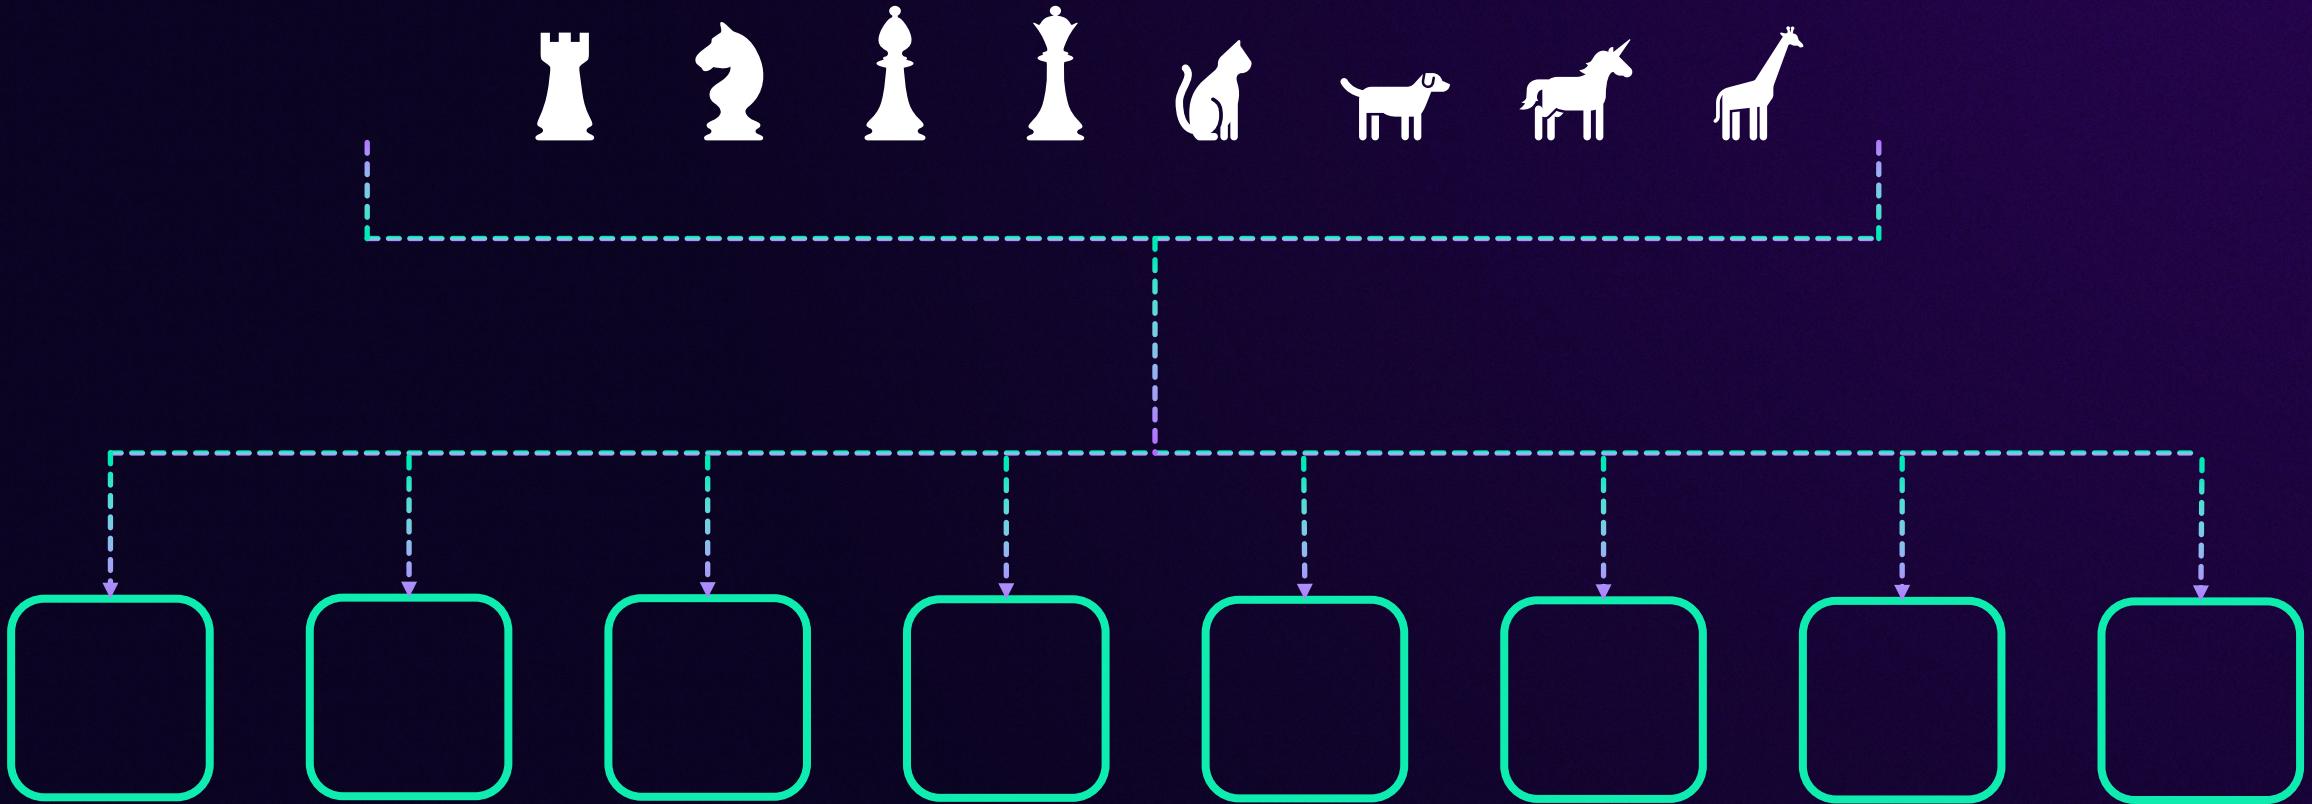


Scope of impact = All customers

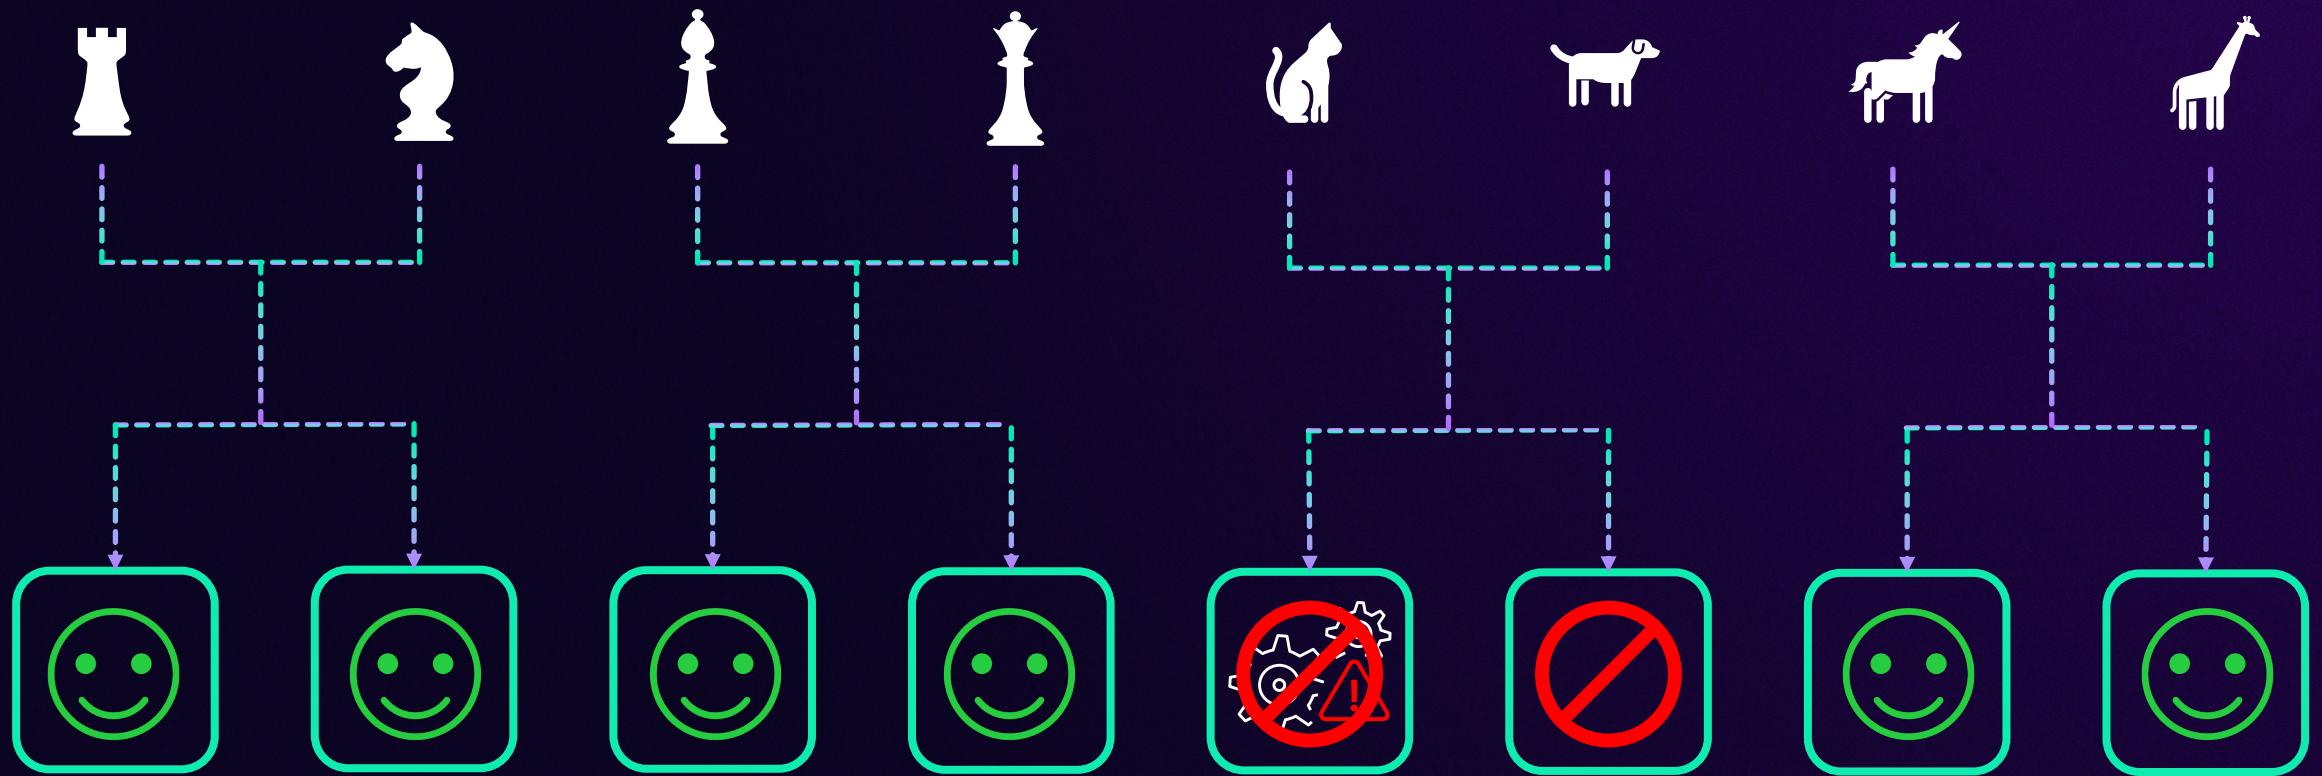
# Throttling



# Sharding



# Sharding

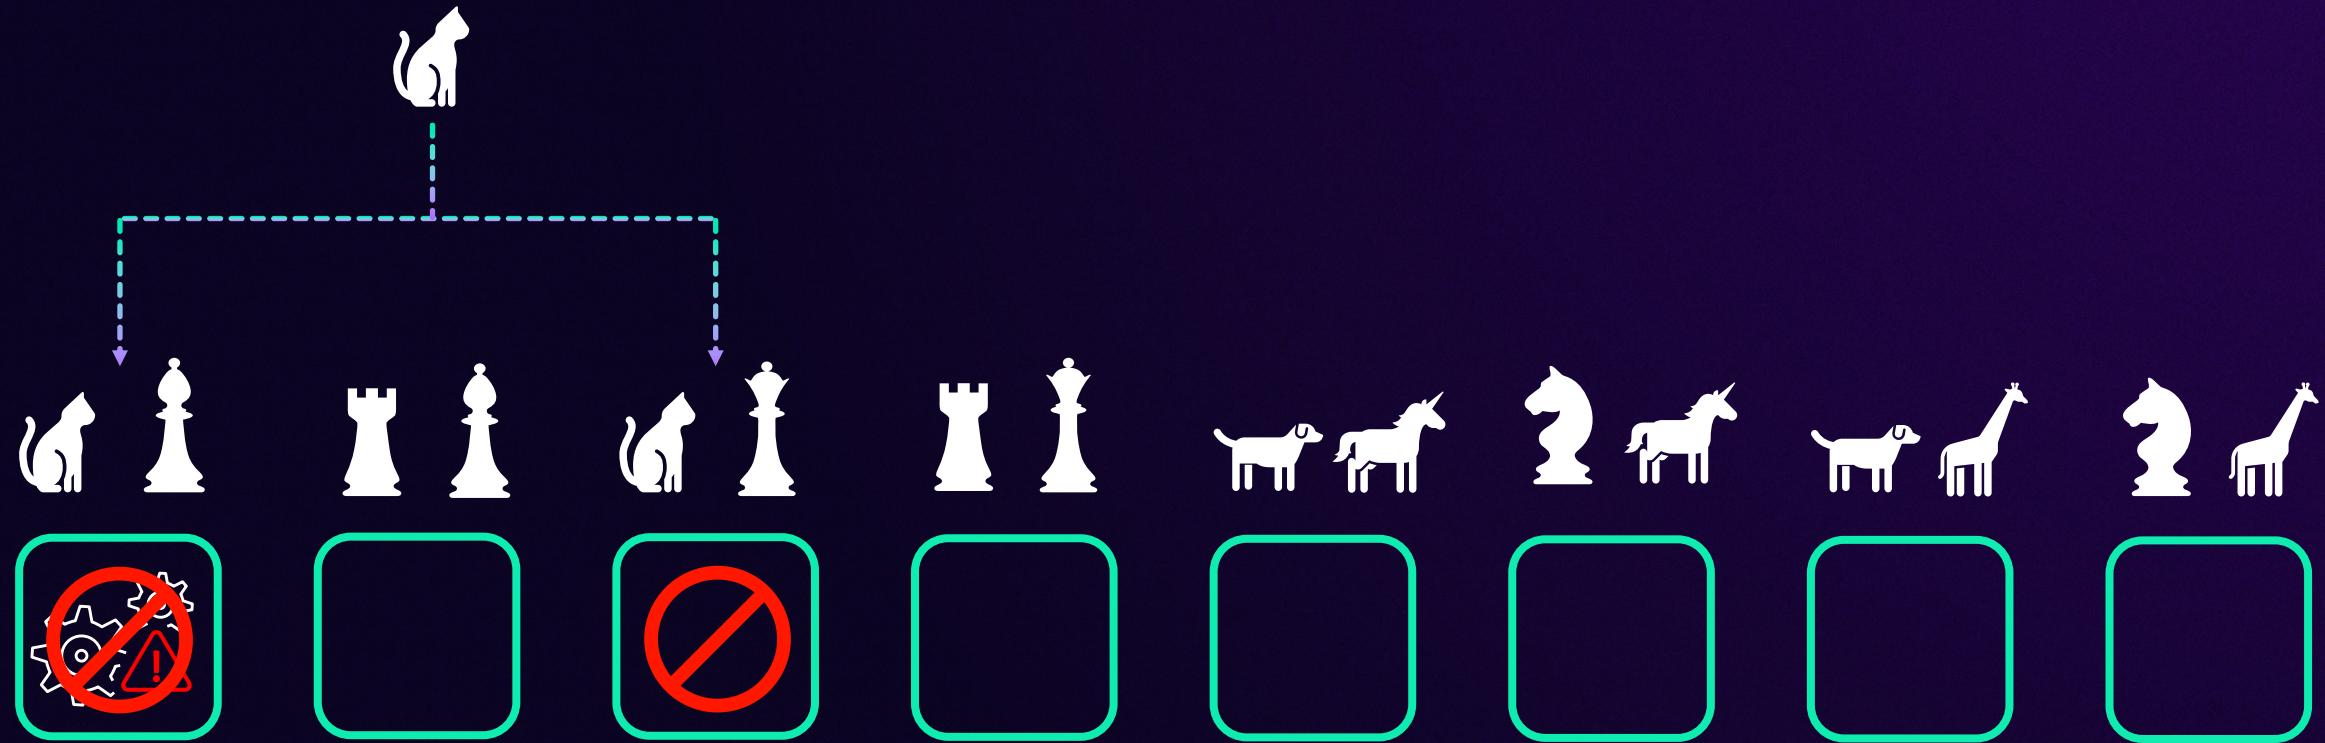


Scope of impact =  $\frac{\text{Customers}}{\text{Shards}}$

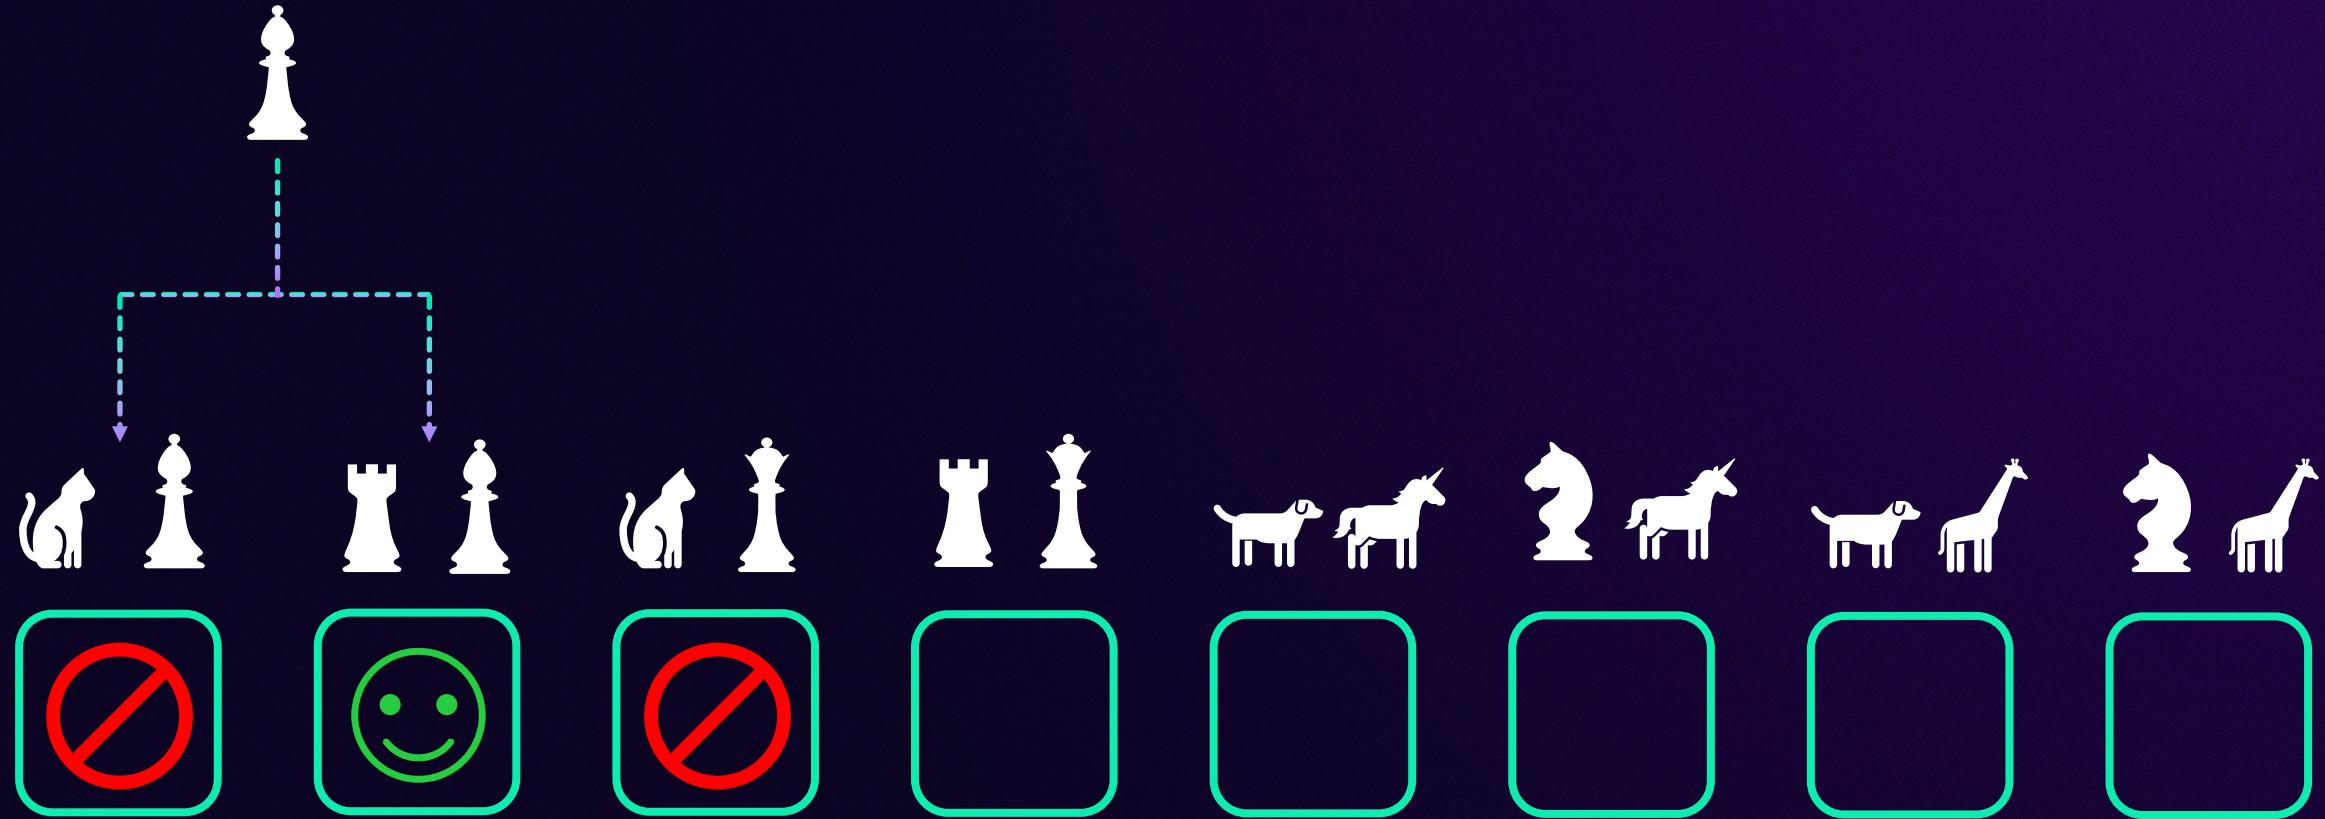
# Shuffle sharding



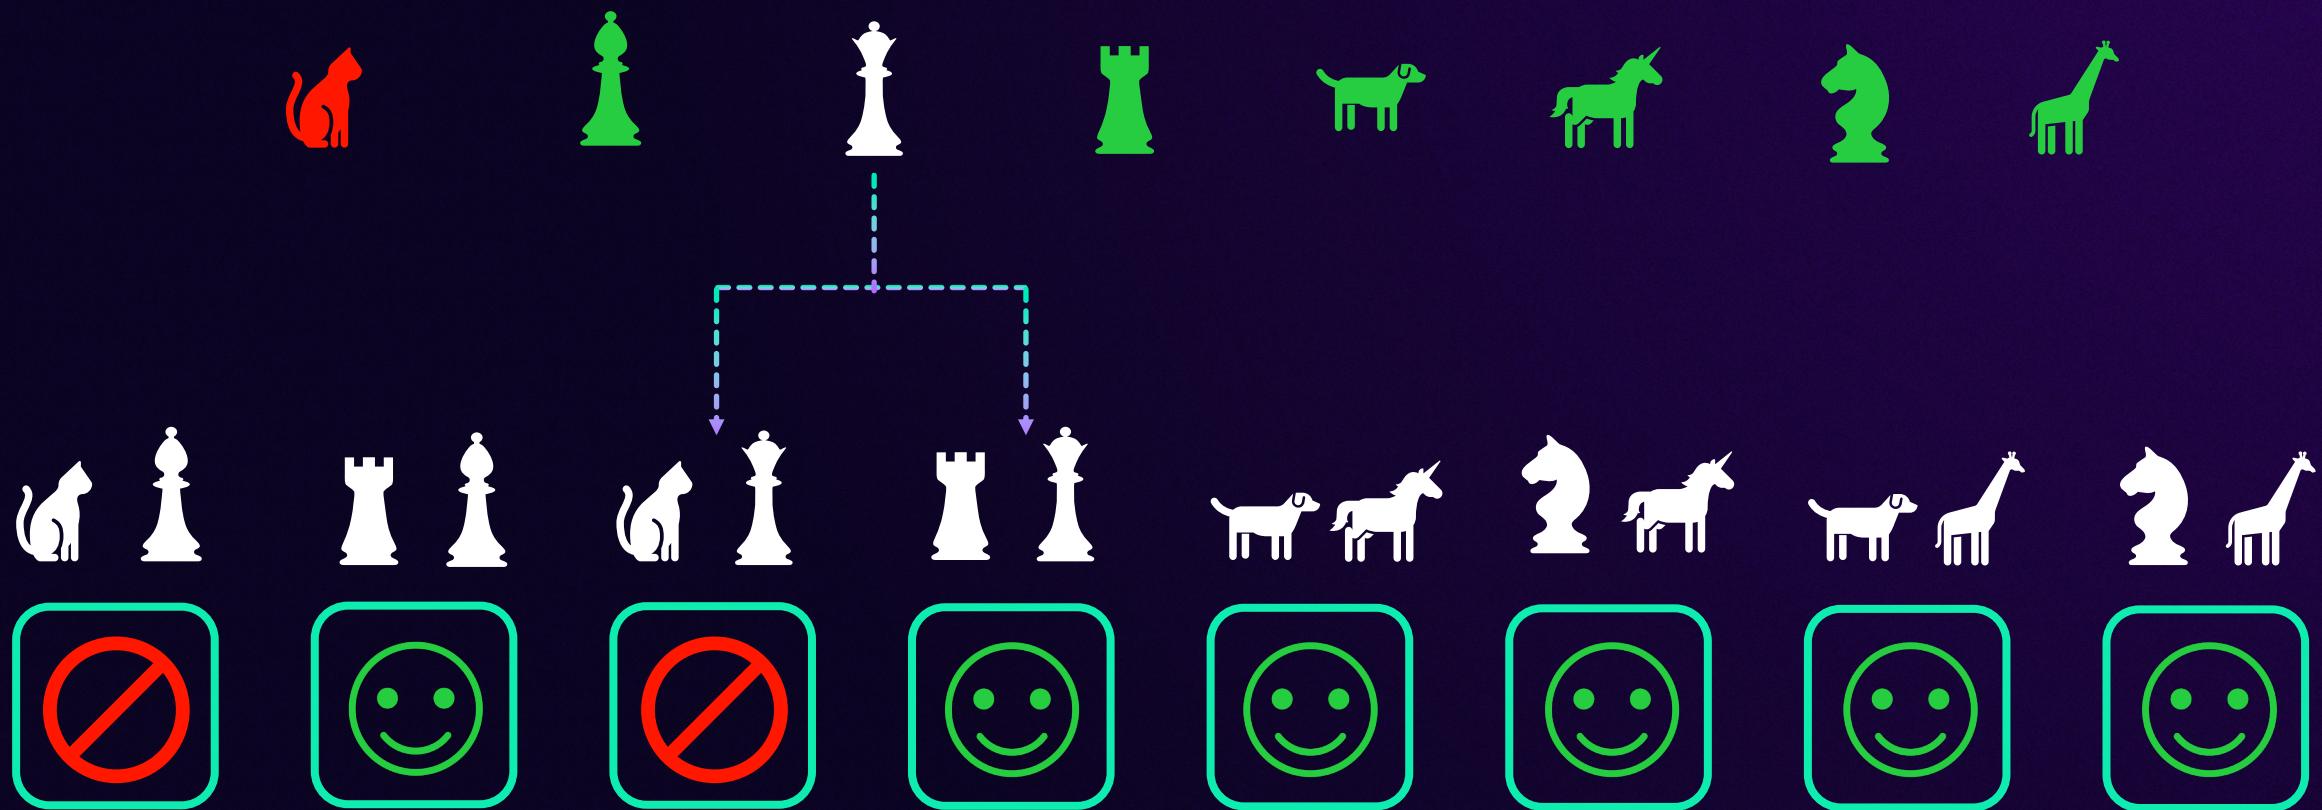
# Shuffle sharding



# Shuffle sharding



# Shuffle sharding



$$\text{Scope of impact} = \frac{\text{Customers}}{\text{Combinations}}$$

# Multi-Tenant (SaaS) architectures

## Challenge



One size fits all mindset

## Tool

Choose the right tenant isolation model and be flexible

---



Tenant identity & isolation



Blast radius & Noisy neighbors

JWT tokens, API-, Service-, Encryption-based Runtime policies

---

Cell-Based architectures, Throttling and Shuffle-Sharding

# Thank you!

**Maximilian Schellhorn**

[mxschell@amazon.com](mailto:mxschell@amazon.com)

[linkedin.com/in/maxschell/](https://linkedin.com/in/maxschell/)



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved.